

UNIVERSITY OF SÃO PAULO
INSTITUTE OF MATHEMATICS AND STATISTICS
BACHELOR OF COMPUTER SCIENCE

**Argumentation mining using
neuro-probabilistic answer set programs**

Jonas Rodrigues Lima Gonçalves

FINAL ESSAY
MAC 499 — CAPSTONE PROJECT

Supervisor: Prof. Dr. Denis Deratani Mauá

São Paulo
2025

I authorize the complete or partial reproduction and disclosure of this work by any conventional or electronic means for study and research purposes, provided that the source is acknowledged.

Ficha catalográfica elaborada com dados inseridos pelo(a) autor(a)
Biblioteca Carlos Benjamin de Lyra
Instituto de Matemática e Estatística
Universidade de São Paulo

Lima Gonçalves, Jonas Rodrigues
Argumentation mining using neuro-probabilistic answer
set programs / Jonas Rodrigues Lima Gonçalves; orientador,
Denis Maua. - São Paulo, 2025.
98 p.: il.

Tese de Conclusão de Curso - Programa de Graduação
em Ciência da Computação / Instituto de Matemática e
Estatística / Universidade de São Paulo.
Bibliografia
Versão original

1. Inteligência Artificial. 2. Mineração de Argumentos.
3. Programação Lógico Probabilística. 4. Compilação
de Conhecimento. 5. Circuitos Probabilísticos. I.
Maua, Denis. II. Título.

Bibliotecárias do Serviço de Informação e Biblioteca
Carlos Benjamin de Lyra do IME-USP, responsáveis pela
estrutura de catalogação da publicação de acordo com a AACR2:
Maria Lúcia Ribeiro CRB-8/2766; Stela do Nascimento Madruga CRB 8/7534.

Acknowledgements

I am profoundly grateful to my advisor, Professor Denis Deratani Mauá, for his invaluable guidance and insightful feedback throughout this work. His mentorship has been instrumental in shaping my academic trajectory and fostering my intellectual growth. I extend my heartfelt thanks to my family, whose unwavering encouragement, patience, and belief in me have been a constant source of strength. I am also deeply appreciative of my friends and colleagues at IME-USP for their camaraderie, stimulating discussions, and the inspiration they have provided along the way. All this support has made this journey not only rewarding but also truly memorable.

Abstract

Jonas Rodrigues Lima Gonçalves. **Argumentation mining using neuro-probabilistic answer set programs.** Capstone Project Report (Bachelor). Institute of Mathematics and Statistics, University of São Paulo, São Paulo, 2025.

Argumentation mining is a complex task that has been approached more recently using purely connectionist methods. These methods aim to extract arguments from text data and represent them in a structured format. However, these methods often lack the ability to handle uncertainty and probabilistic reasoning, and usually require large amounts of labeled data for training, which often are not available in many real-world scenarios.

In order to address these limitations, frameworks for modeling and reasoning about arguments have been developed. These frameworks model the problem via a Neuro-Symbolic approach, either using Integer Linear Programming or Probabilistic Logic Programming (PLP). While Integer Linear Programming is a well-known mathematical optimization technique that can be used to model and solve complex decision-making problems, the integration of such methods to constrain learning and inference is a challenging task. On the other hand, Probabilistic Logic Programming stands itself as a powerful tool with declarative semantics and more “readable” syntax for non-experts.

As of the time of writing, current PLP frameworks for modeling Argumentation Mining focused on using stratified programs, which largely restrict the expressiveness of the different argumentation problems one may desire to represent. Thus, we propose to model this problem using Probabilistic Answer Set Programming (PASP), a framework that combines the expressiveness of ASP with the probabilistic reasoning capabilities of PLP. Furthermore, in order to be able to use this PASP framework for scalable Neuro-Symbolic learning, we explore different state-of-the-art Knowledge Compilation (KC) techniques of the language, which are able to encode PASP programs into circuits that can be encoded as computational graphs in a variety of autodifferentiable frameworks, such as PyTorch or Jax.

Keywords: Probabilistic Logic Programming. Knowledge Compilation. Probabilistic Circuits. Argument Mining.

Resumo

Jonas Rodrigues Lima Gonçalves. **Mineração de argumentos usando programas de conjuntos de respostas probabilísticas.** Monografia (Bacharelado). Instituto de Matemática e Estatística, Universidade de São Paulo, São Paulo, 2025.

A mineração de argumentação é uma tarefa complexa que tem sido abordada mais recentemente utilizando métodos puramente conexionistas. Esses métodos têm como objetivo extrair argumentos de dados textuais e representá-los em um formato estruturado. No entanto, esses métodos frequentemente carecem da capacidade de lidar com incertezas e raciocínio probabilístico, além de geralmente exigirem grandes quantidades de dados rotulados para treinamento, os quais muitas vezes não estão disponíveis em muitos cenários do mundo real.

Para lidar com essas limitações, foram desenvolvidos frameworks para modelar e realizar raciocínio sobre argumentos. Esses frameworks modelam o problema por meio de uma abordagem Neuro-Simbólica, utilizando Programação Linear Inteira ou Programação Lógica Probabilística (PLP). Enquanto a Programação Linear Inteira é uma técnica de otimização matemática bem conhecida que pode ser usada para modelar e resolver problemas complexos de tomada de decisão, a integração de tais métodos para restringir aprendizado e inferência é uma tarefa desafiadora. Por outro lado, a Programação Lógica Probabilística se destaca como uma ferramenta poderosa com semântica declarativa e uma sintaxe mais “legível” para não especialistas.

No momento da redação deste trabalho, os frameworks atuais de PLP para modelar a Mineração de Argumentação focaram no uso de programas estratificados, o que restringe amplamente a expressividade dos diferentes problemas de argumentação que se deseja representar. Assim, propomos modelar esse problema utilizando a Programação Lógica com Conjuntos de Respostas Probabilística (PASP), um framework que combina a expressividade da ASP com as capacidades de raciocínio probabilístico da PLP. Além disso, para ser capaz de usar esse framework PASP para aprendizado Neuro-Simbólico escalável, exploramos diferentes técnicas modernas de Compilação de Conhecimento (KC) da linguagem, que são capazes de codificar programas PASP em circuitos que podem ser representados como grafos computacionais em uma variedade de frameworks autodiferenciáveis, como PyTorch ou Jax.

Palavras-chave: Programação Lógico Probabilística. Compilação de Conhecimento. Circuito Probabilísticos. Mineração de Argumentos.

List of Abbreviations

2AMC	Second Level Algebraic Model Counting
AASC	Algebraic Answer Set Counting
AC	Arithmetic Circuit
AMC	Algebraic Model Counting
ASC	Answer Set Counting
BDD	Binary Decision Diagram
DNNF	Decomposable Negation Normal Form
d-DNNF	Deterministic Decomposable Negation Normal Form
EDLP	Extended Disjunctive Logic Programming
ELP	Extended Logic Programming
FOL	First Order Logic
KC	Knowledge Compilation
KR	Knowledge Representation
LC	Logic Circuit
LP	Logic Programming
NNF	Negation Normal Form
NP	Nondeterministic Polynomial Time
OBDD	Ordered Binary Decision Diagram
PASP	Probabilistic Answer Set Programming
PC	Probabilistic Circuit
PLP	Probabilistic Logic Programming
PSDD	Probabilistic Sentential Decision Diagram
SDD	Sentential Decision Diagram
sd-DNNF	Smooth Decomposable Negation Normal Form
SDNNF	Structured Decomposable Negation Normal Form
WMC	Weighted Model Counting

List of Symbols

X, Y, Z, \dots	Sets of variables
x, y, z, \dots	Sets of assignments
$\langle f \rangle$	Semantics of Boolean formula f
$f \equiv g$	Equivalence between Boolean formulae f and g (i.e. $\langle f \rangle = \langle g \rangle$)
N, S, P, L	Sets of nodes
$Ch(N)$	Set of all children of node N
$Pa(N)$	Set of all parents of node N
$Desc(N)$	Set of all descendants of node N
$head(r)$	Head of a rule r
$body(r)$	The set of atoms inside the body of a rule r

List of Figures

5.1	Knowledge compilation pipeline with grounding, compilation, and inference stages.	29
5.2	Illustrative bipolar argumentation fragment with mutual support and derived defeat chains.	30
5.3	Smooth, decomposable, and deterministic circuit fragment used for PASP inference.	31

List of Tables

- | | |
|---|---|
| 3.1 Examples of logical and probabilistic inference tasks that can be modelled as AMC instances, together with their semirings and labelling functions. | 9 |
|---|---|

List of Programs

Contents

1	Introduction	1
2	Appendix Satisfiability and Proof Theory	3
2.1	Notation	3
2.2	Satisfiability	4
2.2.1	Satisfiability	5
2.2.2	MAX-SAT	5
2.2.3	Sharp-SAT	6
3	Probabilistic Logical Inference	7
3.1	From SAT to AMC	7
3.2	Semirings	7
3.3	Algebraic Model Counting	8
3.4	Second Level of Algebraic Model Counting	10
4	Answer Set Programming	13
4.1	Foundations of Logic Programming	13
4.1.1	Definite Logic Programs	13
4.1.2	Normal Logic Programs	14
4.1.3	Herbrand Universe, Base and Interpretation	14
4.1.4	Grounding of a Logic Program	15
4.1.5	Interpretation of a Logic Program	15
4.1.6	Loop Formulas	16
4.1.7	Immediate Consequence Operator	16
4.1.8	Negation and Stratification	17
4.1.9	Stable Model Semantics	18
4.1.10	Reasoning	19
4.2	Answer Set Programming	19
4.2.1	Extended Logic Programs	19

4.2.2	Cardinality Constraints	21
4.3	Probabilistic Logic Programming	21
4.3.1	Probabilistic Answer Set Programming	22
4.3.2	PASP Inference as 2AMC	25
5	Contribution	27
5.1	Scope and Roadmap	27
5.2	Problem Setting	27
5.2.1	From Argumentation Mining to PASP	27
5.2.2	Knowledge Compilation Pipeline	28
5.3	Foundational Building Blocks	28
5.3.1	Argumentation Frameworks	28
5.3.2	Knowledge Compilation	29
5.4	Experimental Results	31
5.4.1	Experimental Outlook	31
5.4.2	Expected Outcomes	33
5.4.3	Empirical Analysis	33
References		35

Chapter 1

Introduction

Argumentation is a fundamental aspect of human communication, surfacing in contexts that range from informal discussions to legal reasoning. Beyond persuasion, argumentative exchanges support critical thinking, collective decision-making, and the construction of shared knowledge. Understanding how arguments are structured, how they interact, and how they can be evaluated is therefore a central challenge for both the social sciences and Artificial Intelligence (AI).

Recent advances in argumentation mining attempt to automatically detect the building blocks of argumentative discourse (claims, premises, supports, and attacks) directly from text (STAB and GUREVYCH, 2017). The quality of these pipelines depends on coordinating several interdependent subtasks: deciding whether a span of text is argumentative, classifying its role, and predicting how it relates to the rest of the discourse. Errors in the early stages easily propagate to later ones, eroding the coherence of the recovered argumentative graph. Neuro-symbolic methods are especially attractive in this setting because they combine neural networks, which excel at processing unstructured language, with symbolic reasoning, which enforces global consistency and domain constraints.

The foundational joint model of STAB and GUREVYCH (2017) couples local classifiers with Integer Linear Programming (ILP) constraints that capture the shape of well-formed argumentative structures. ILP-based approaches improved consistency over pipeline baselines, yet they struggle to represent uncertainty and incur a combinatorial cost when the search space grows. Probabilistic Logic Programming (PLP) systems, such as ProbLog (FIERENS *et al.*, 2015), mitigate these limitations by natively handling stochastic information while retaining the declarative and explainable nature of logic programs. Frameworks like DeepProbLog (MANHAEVE *et al.*, 2018) and its application to argumentation mining (CERVEIRA DO AMARAL *et al.*, 2023) demonstrate how neural predictions can be wired into symbolic models that reason about argumentative structure.

Extending this line of work, sMPROBLOG adopts the stable model semantics of Answer Set Programming (ASP), unlocking the ability to represent negative cycles and other non-monotonic phenomena that occur in real-world argumentation (TOTIS *et al.*, 2023). Stable models are expressive enough to encode the bipolar argumentation frameworks we study in this thesis, but their probabilistic extensions quickly lead to intractable inference

when treated naïvely.

This dissertation examines how probabilistic ASP (PASP) can be compiled into structured probabilistic circuits that support tractable inference and learning. By embedding argumentative constraints into compiled circuits we seek to align neural predictions with the semantics of bipolar argumentation, enabling transparent reasoning, calibrated uncertainty, and scalable training loops. The central questions we explore are:

- How can argumentation mining tasks be encoded as PASP programs that capture the mutual influence of claims, premises, supports, and attacks while preserving probabilistic semantics for uncertain or conflicting evidence?
- Which knowledge compilation strategies can translate these PASP encodings into circuits that admit efficient marginal and conditional inference?
- How can the resulting circuits interface with neuro-symbolic learning, allowing gradients or other learning signals to flow through compiled structure without sacrificing explainability?

The remainder of this document is organized as follows. Chapters [chapter 2](#), [chapter 3](#), and [chapter 4](#) review background on propositional satisfiability, algebraic model counting, and logic programming with stable models. Chapter [chapter 5](#) then introduces our PASP encoding of bipolar argumentation, the knowledge compilation pipeline that supports it, and the circuit properties required for efficient inference. We later position our approach with respect to the literature and discuss the experimental questions that guide our evaluation.

Chapter 2

Appendix Satisfiability and Proof Theory

Logic revolves around two tightly linked notions: *consistency* and *validity* (FRANCO and MARTIN, 2009). Proof theory studies these notions from a syntactic perspective by manipulating symbols using axioms and inference rules; model theory examines them semantically by interpreting sentences over structures that encode the “world”. Although the two views emphasise different objects, they coincide for complete axiom systems: a statement that can be derived syntactically is precisely one that is true in every model, and unsatisfiable theories are those that entail a contradiction.

The bridge between proof theory and model theory is central to this thesis. Efficient algorithms for propositional satisfiability (SAT) allow us to reason about the models of logic programs, while proofs about circuit properties ensure that the compiled artefacts faithfully capture the semantics of the original theories. Shannon’s connection between Boolean logic and circuits (SHANNON, 1938) sparked a line of work that culminated in modern knowledge compilation techniques, such as Binary Decision Diagrams (BDDs) (AKERS, 1978; BRYANT, 1986) and Sentential Decision Diagrams (SDDs). These structures will later support our probabilistic Answer Set Programming encodings.

2.1 Notation

We borrow definitions and notation from the work of (ARORA and BARAK, 2009) and (KIMMIG *et al.*, 2017). Since this chapter is focused on logical SAT-like problems, we only introduce basic logic concepts, such as *Boolean formulas*, CNFs and *Logical Propositional Theories*.

Definition 2.1.1 (Boolean Formulas). A *Boolean formula* is the combination of *Boolean variables*, that can be either \top or \perp , and logical operators, such as \wedge (*AND*), \vee (*OR*), \neg (*NOT*, also denoted by an overline) (ARORA and BARAK, 2009).

Example 2.1.1. For example, the formula ϕ defined by

$$\phi = (u_1 \wedge u_2) \vee \neg(u_3 \wedge \bar{u}_4)$$

is a Boolean formula over variables u_1, u_2, u_3, u_4 . We denote by $\phi(a)$ the evaluation of ϕ over an assignment a of the variables in the formula.

Definition 2.1.2 (Conjunctive Normal Form). A Boolean formula over variables u_1, \dots, u_n is in CNF if it is a conjunction of clauses, where each clause is a disjunction of literals, and a literal is either a variable or its negation (ARORA and BARAK, 2009). Hence, a CNF is a Boolean formula of the form

$$\bigwedge_i \left(\bigvee_j v_{ij} \right),$$

where each v_{ij} is a literal over variables u_1, \dots, u_n .

Example 2.1.2. For example, the following formula is a 3CNF (a CNF where each clause has at most 3 literals) over variables u_1, u_2, u_3, u_4 :

$$(u_1 \vee \bar{u}_2 \vee u_3) \wedge (u_2 \vee \bar{u}_3 \vee u_4) \wedge (\bar{u}_1 \vee u_3 \vee \bar{u}_4).$$

Definition 2.1.3 (Disjunctive Normal Form). A Boolean formula over variables u_1, \dots, u_n is in DNF if it is a disjunction of clauses, where each clause is a conjunction of literals. Hence, a DNF is a Boolean formula of the form:

$$\bigvee_i \left(\bigwedge_j v_{ij} \right),$$

where each v_{ij} is a literal over variables u_1, \dots, u_n .

Definition 2.1.4 (Logical Propositional Theory). A *Logical Theory* is set of logical sentences, Boolean-valued formulas with no free variables. Hence, a *Propositional Theory* is Logical Theory where the sentences are propositional formulas, i.e., formulas with no quantifiers.

2.2 Satisfiability

Propositional satisfiability (SAT) is one of the cornerstones of theoretical computer science. Every problem in NP polynomially reduces to SAT, which establishes its NP-completeness (COOK, 1971; LEVIN, 1973). Because of this universality, SAT provides a convenient target for reasoning about diverse logical theories. Later chapters will rely on SAT when translating PASP encodings into propositional form so that we can leverage mature SAT and model counting machinery.

In neuro-symbolic settings we seldom solve SAT just once; instead we evaluate many related queries that differ by evidence or parameter values. This calls for solving SAT-like problems in a way that enables reuse, motivating the knowledge compilation techniques

studied throughout this thesis. We therefore start by recalling baseline notions of satisfiability and its counting and optimization variants.

2.2.1 Satisfiability

The Boolean Satisfiability Problem is the problem of determining whether a given CNF is *satisfiable*, i.e., whether there is an assignment of *True* or *False* to the variables that makes the formula *True* (ARORA and BARAK, 2009). If such an assignment does not exist, the formula is said to be *unsatisfiable*.

Definition 2.2.1 (SAT - Satisfiability). Let SAT be the language of all satisfiable CNFs. Then, we say that for a CNF ϕ , the SAT problem is defined as

$$\text{SAT}(\phi) = \{\phi \mid \phi \text{ is a satisfiable propositional formula}\}.$$

For example, $(\text{True}, \text{False}, \text{False}, \text{False})$ is an example of assignment that makes the 3CNF of definition 2.1.2 satisfiable. On the other hand, the CNF $x \wedge \bar{x}$ is unsatisfiable, since it is a *contradiction* (the opposite of a tautology).

2.2.2 MAX-SAT

Max-SAT generalises satisfiability by searching for the assignment that optimises how many clauses evaluate to true (KRENTEL, 1988). The classical (unweighted) variant simply counts satisfied clauses.

Definition 2.2.2 (MAX-SAT). Let ϕ be a CNF whose clause set is C . The Max-SAT objective is

$$\max_{a \in \{0,1\}^n} \sum_{c \in C} \llbracket c(a) \rrbracket,$$

where $\llbracket c(a) \rrbracket$ is the Iverson bracket that yields 1 when clause c evaluates to *True* under assignment a and 0 otherwise, and n is the number of variables in ϕ .

We can additionally attach a non-negative weight $w(c)$ to each clause and seek the maximum weighted sum of satisfied clauses.

Definition 2.2.3 (Weighted MAX-SAT). The weighted Max-SAT problem is

$$\max_{a \in \{0,1\}^n} \sum_{c \in C} w(c) \llbracket c(a) \rrbracket.$$

SAT is a special case of Max-SAT: a CNF is satisfiable iff the optimum equals $|C|$. Consequently, Max-SAT is NP-hard. Setting all weights to 1 reduces the weighted variant to the unweighted one, confirming that the weighted problem inherits NP-hardness. Krentel further shows that Max-SAT is OptP-complete, placing it alongside optimisation problems such as TSP and KNAPSACK (KRENTEL, 1988).

2.2.3 Sharp-SAT

Unlike SAT and Max-SAT, the $\#SAT$ (Sharp-SAT) problem is a counting task: it asks how many satisfying assignments a CNF admits ([VALIANT, 1979](#)). This places $\#SAT$ closer to the counting problems studied by Levin ([LEVIN, 1973](#)) than to Cook's decision formulation ([COOK, 1971](#)): there is always an answer, but enumerating it may require exploring all assignments rather than stopping after the first model is discovered.

Counting solutions is at least as hard as deciding their existence. Every SAT instance can be solved by verifying whether the corresponding $\#SAT$ count is non-zero, so $\#SAT$ is NP-hard. Valiant further proved that it is $\#P$ -complete, meaning that any counting problem associated with an NP decision problem can be reduced to $\#SAT$.

Definition 2.2.4 ($\#SAT$). Formally, let ϕ be a CNF over variables x_1, \dots, x_n . The Sharp-SAT problem asks for the quantity

$$\#SAT(\phi) = |\{a \in \{0, 1\}^n \mid \phi(a) = \text{True}\}|.$$

Similar to the Max-SAT problem, there is a weighted version of the Sharp-SAT problem, called Weighted Model Counting (WMC), where a non-negative weight is associated to each assignment of values to variables ([CHAKRABORTY *et al.*, 2015](#)). The classical version of the problem, usually called Model Counting (MC), can be seen as a special case of the weighted version, where all weights are equal to 1, and is sometimes referred to as *unweighted model counting*.

Definition 2.2.5 (Weighted Model Counting). The Weighted Model Counting problem is the problem of finding the sum of the weights of all assignments that satisfy a given CNF. Formally, let ϕ be a CNF over variables x_1, \dots, x_n , and let $w : \{0, 1\}^n \rightarrow \mathbb{R}^+$ be the weight function over the assignments of the variables in ϕ . Then, the WMC problem is

$$\text{WMC}(\phi) = \sum_{a \in \{0, 1\}^n} w(a) \cdot [\phi(a)].$$

Weighted model counting subsumes the previous variants: setting all weights to 1 recovers model counting, while restricting weights to $\{0, 1\}$ captures ordinary SAT. In Chapter [chapter 3](#) we will further generalise this view using algebraic structures that allow us to express probabilistic inference and other reasoning tasks within a single framework.

Chapter 3

Probabilistic Logical Inference

Chapter chapter 2 introduced SAT and its counting and optimisation variants. Weighted model counting (WMC) is especially important for us because it reduces many probabilistic inference tasks to summing weights over satisfying assignments (Adnan DARWICHE, 2002; CHAVIRA and Adnan DARWICHE, 2008). By encoding, for example, a Bayesian network as a propositional theory whose literals carry probabilities, WMC answers marginal queries through purely logical machinery.

Knowledge compilation builds on this reduction: once a theory has been compiled into a circuit that supports fast WMC, evidence and query evaluation become linear in the circuit size (CHAVIRA, Adnan DARWICHE, and JAEGER, 2006; KIMMIG *et al.*, 2017). To reason about PASP we need a framework that encompasses classical SAT, WMC, and other algebraic tasks. Algebraic Model Counting (AMC) provides precisely that unifying perspective (KIMMIG *et al.*, 2017). A recent extension, second level AMC (2AMC), further captures probabilistic ASP inference (KIESEL, TOTIS, *et al.*, 2022). This chapter introduces these abstractions and prepares the ground for the compilation pipeline described later.

3.1 From SAT to AMC

Before defining AMC we recall the algebraic structure that underpins it. Generalising from simple numeric addition and multiplication allows us to model counting, optimisation, and probability computations within the same template.

3.2 Semirings

Definition 3.2.1 (Semiring). A semiring is an algebraic structure $(\mathcal{A}, \oplus, \otimes, e_{\oplus}, e_{\otimes})$ where addition \oplus and multiplication \otimes are associative binary operations on \mathcal{A} , \oplus is commutative, \otimes distributes over \oplus , e_{\oplus} and e_{\otimes} are neutral elements for \oplus and \otimes , respectively, and e_{\oplus} annihilates multiplication: $e_{\oplus} \otimes a = a \otimes e_{\oplus} = e_{\oplus}$ for every $a \in \mathcal{A}$. In a *commutative* semiring, multiplication \otimes is also commutative.

Semirings appear throughout computer science. The *tropical* semiring, for instance,

replaces addition with the minimum operator and multiplication with the standard addition over extended real numbers (SIMON, 1988). When used with AMC it captures optimisation tasks like Max-SAT: summations aggregate clause scores while products accumulate costs along a derivation.

3.3 Algebraic Model Counting

Definition 3.3.1 (Algebraic Model Counting). Given

- A propositional logic theory T over a set of variables \mathcal{V} ;
- A commutative semiring $(\mathcal{A}, \oplus, \otimes, e_{\oplus}, e_{\otimes})$; and
- A labelling function $\alpha : \mathcal{L} \rightarrow \mathcal{A}$, mapping literals \mathcal{L} of the variables in \mathcal{V} to elements of the semiring set \mathcal{A} .

The AMC problem is now defined as the computation of the following expression:

$$A(T) = \bigoplus_{I \in \mathcal{M}(T)} \bigotimes_{i \in I} \alpha(i),$$

where $\mathcal{M}(T)$ denotes the set of models of T .

To anyone familiar with introductory algebra, it is easy to see that the construction above subsumes many SAT-like problems. For example, by setting the semiring to

$$(\mathcal{A}, \oplus, \otimes, e_{\oplus}, e_{\otimes}) = (\{\text{true}, \text{false}\}, \vee, \wedge, \text{false}, \text{true}),$$

and mapping every literal to *true* through α , the AMC expression collapses to Boolean reasoning and therefore decides SAT. Moreover, by choosing

$$(\mathcal{A}, \oplus, \otimes, e_{\oplus}, e_{\otimes}) = (\mathbb{N}, +, \times, 0, 1),$$

and mapping both positive and negative literals to 1, the resulting AMC computation counts how many models satisfy the theory.

Similarly, by only changing the set \mathcal{A} of the semiring to $\mathbb{R}_{\geq 0}$ (non-negative real numbers) and letting α assign literal weights in $\mathbb{R}_{\geq 0}$, the AMC problem recovers weighted model counting.

Moreover, the results cited previously about Bayesian networks imply that AMC also captures probabilistic inference (ADNAN DARWICHE, 2002; CHAVIRA, ADNAN DARWICHE, AND JAEGER, 2006; CHAVIRA AND ADNAN DARWICHE, 2008; SANG *et al.*, 2005). The algebraic structure makes the connection explicit: using the same semiring and defining $\alpha(v) \in [0, 1]$ as the probability of literal v , with $\alpha(\neg v) = 1 - \alpha(v)$ for binary variables, yields the desired probability mass assignments.

Example 3.3.1. Consider the theory $T = (a \vee b)$ and the probability semiring $(\mathbb{R}_{\geq 0}, +, \times, 0, 1)$. Let the labelling function assign weights $\alpha(a) = 0.6$, $\alpha(\neg a) = 0.4$, $\alpha(b) = 0.3$, and $\alpha(\neg b) = 0.7$.

The models that satisfy T are all assignments except $(\neg a, \neg b)$. Applying the AMC formula yields

$$A(T) = 1 - \alpha(\neg a)\alpha(\neg b) = 1 - 0.4 \times 0.7 = 1 - 0.28 = 0.72,$$

which coincides with the probability that at least one of a or b is true. This simple computation mirrors the WMC task we will perform on compiled circuits.

In fact, AMC is capable of modeling many other interesting problems, such as *EXPEC* (expectation), which allows one to infer parameters in a finite-state transducer relative to a given dataset. The elements of the respective expectation semiring are tuples of the form (p, v) , where $p \in [0, 1]$ is the probability of traversing a particular arc and $v \in \mathbb{R}$ is the value contributed by that arc. The operations \oplus, \otimes , neutral elements e_{\oplus}, e_{\otimes} , and a suitable labelling function are defined as follows:

$$\begin{aligned} (p_1, v_1) \oplus (p_2, v_2) &= (p_1 + p_2, v_1 + v_2), \\ (p_1, v_1) \otimes (p_2, v_2) &= (p_1 \cdot p_2, p_1 \cdot v_2 + p_2 \cdot v_1), \\ e_{\oplus} &= (0, 0), \\ e_{\otimes} &= (1, 0), \\ \alpha(a) &= (p_a, p_a \cdot v_a), \end{aligned}$$

where each arc a carries probability p_a and reward v_a .

Adapted from (KIMMIG *et al.*, 2017), Table 3.1 summarises AMC instances for different logical and probabilistic tasks, detailing the underlying sets, operations, and labelling functions for each semiring.

Task	\mathcal{A}	\oplus	\otimes	e^{\oplus}	e^{\otimes}	$\alpha(v)$	$\alpha(\neg v)$
SAT	$\{\text{true}, \text{false}\}$	\vee	\wedge	false	true	true	true
#SAT	\mathbb{N}	$+$	\times	0	1	1	1
WMC	$\mathbb{R}_{\geq 0}$	$+$	\times	0	1	$\in \mathbb{R}^+$	$\in \mathbb{R}^+$
PI	$\mathbb{R}_{\geq 0}$	$+$	\times	0	1	$\in [0, 1]$	$1 - \alpha(v)$

Table 3.1: Examples of logical and probabilistic inference tasks that can be modelled as AMC instances, together with their semirings and labelling functions.

The semirings associated with SAT and probabilistic inference are often called the *Boolean* and the *probability* semirings, respectively.

Other useful problems that can be modeled by the AMC task are: *sensitivity analysis* (SENS), *probability of most likely states* (MPE), *shortest* and *widest path* (S-PATH and W-PATH, respectively), *fuzzy* and *k-weighted* constraints (FUZZY and k-WEIGHT, respectively), and *OBDD_C construction* (KIMMIG *et al.*, 2017).

3.4 Second Level of Algebraic Model Counting

A natural generalisation of AMC introduces a third operation alongside \oplus and \otimes . The resulting framework, called 2AMC, captures tasks such as Maximum a Posteriori (MAP) inference and is essential when modelling PASP problems, as we will see in Chapter 4.

We now recall the definition of the 2AMC problem (KIESEL, TOTIS, *et al.*, 2022):

Definition 3.4.1 (Second-Level Algebraic Model Counting). Given

- A propositional logic theory T over a set of variables \mathcal{V} ;
- A partition of the variables in T , (X_I, X_O) ;
- Two commutative semirings $S_j = (\mathcal{A}_j, \oplus_j, \otimes_j, e_{\oplus_j}, e_{\otimes_j})$, for $j \in \{I, O\}$;
- Two labelling functions $\alpha_j : \mathcal{L}(X_j) \rightarrow \mathcal{A}_j$, for $j \in \{I, O\}$, mapping literals over the variables in X_j to elements of the semiring set \mathcal{A}_j ; and
- A weight transformation function $t : \mathcal{A}_I \rightarrow \mathcal{A}_O$ that respects $t(e_{\oplus_I}) = e_{\oplus_O}$.

Then the 2AMC problem is defined as the computation of the following expression:

$$2AMC(T) = \bigoplus_{\mathbf{a} \in \mathcal{A}(X_O)}^O \bigotimes_{a \in \mathbf{a}}^O \alpha_O(a) \otimes_O t \left(\bigoplus_{I \in \mathcal{M}(T|\mathbf{a})}^I \bigotimes_{i \in I}^I \alpha_I(i) \right),$$

where $\mathcal{A}(X)$ denotes the set of assignments to the variables in X , and $\mathcal{M}(T|\mathbf{a})$ denotes the set of models of T consistent with assignment \mathbf{a} .

It is easy to see that AMC is a special case of 2AMC. By choosing an empty outer partition ($X_O = \emptyset$) and letting the weight transformation be the identity, the outer summation collapses and we recover the original AMC expression. Intuitively, 2AMC first enumerates assignments to the *outer* variables and re-weights the result of an *inner* AMC computation on the remaining variables before aggregating the outcomes.

A canonical example is Maximum a Posteriori (MAP) inference in probabilistic models. Let T be a propositional theory over variables \mathcal{V} , partition its atoms into queries Q , evidence E , and the remaining hidden variables R . Given evidence e , MAP seeks the most likely assignment to Q :

$$MAP(Q | e) = \arg \max_q \sum_{\mathbf{r} \in \mathcal{A}(R)} P(Q = q, E = e, R = \mathbf{r}),$$

where $\mathcal{A}(R)$ enumerates assignments to R . Within 2AMC we set $X_O = Q$ and $X_I = R \cup E$. The inner AMC sums over the hidden variables for a fixed choice of q and e , while the outer semiring combines these values with a maximisation operator that selects the best query assignment. The weight transformation t merely injects the inner probability mass into the outer semiring so that maximisation can be performed.

To complete the specification we instantiate the two semirings and the weight transformation. The inner semiring is the standard probability semiring $S_I = ([0, 1], +, \times, 0, 1)$ with a labelling function α_I that assigns unit weight to evidence literals consistent with e and

uses the distributional parameters of T for the remaining atoms. Inconsistent literals receive weight 0, ensuring that only models compatible with the evidence contribute to the sum.

The outer semiring performs maximisation. We use $S_O = (\mathbb{R}_{\geq 0} \times 2^{|Q|}, \oplus_{\text{argmax}}, \otimes_{\text{argmax}}, (0, \emptyset), (1, \emptyset))$, where

$$(p_1, q_1) \oplus_{\text{argmax}} (p_2, q_2) = \begin{cases} (p_1, q_1) & \text{if } p_1 > p_2, \\ (p_2, q_2) & \text{if } p_2 > p_1, \\ (p_1, \min\{q_1, q_2\}) & \text{otherwise,} \end{cases}$$

breaking ties by a fixed lexicographic order on assignments, and $(p_1, q_1) \otimes_{\text{argmax}} (p_2, q_2) = (p_1 \cdot p_2, q_1 \cup q_2)$. The labelling α_O associates each literal with its local probability and the singleton set containing that literal. Finally, the weight transformation $t(p) = (p, \emptyset)$ carries the inner probability mass into the outer semiring so that the maximisation can compare different assignments. With these choices, 2AMC reproduces MAP inference exactly.

Chapter 4

Answer Set Programming

Knowledge Representation (KR) research is driven by a familiar tension: expressive formalisms capture rich domains but typically incur intractable inference, while restricted languages admit efficient reasoning at the cost of what can be expressed. Full first-order logic, for example, allows us to encode complex knowledge bases but entailment in this setting is undecidable in general and, even for decidable fragments, often demands exponential resources (LEVESQUE, 1986). Despite numerous refinements, classical resolution-based procedures inherit this worst-case behaviour.

One way to regain tractability is to move to less expressive yet carefully designed fragments. Propositional logic exemplifies this strategy: by grounding away quantifiers we can apply SAT technology, as discussed in Chapter 2. Answer Set Programming (ASP) follows the same philosophy. It restricts syntax in a way that keeps many modelling conveniences—default negation, recursion, and non-monotonic reasoning—while providing a semantics that supports efficient solver implementations. ASP has become a workhorse for applications ranging from planning to computational argumentation (TONI and SERGOT, 2011).

This chapter reviews the logical foundations required for the probabilistic ASP encodings that we adopt later. We begin with definite programs and the associated semantic machinery, introduce stable-model semantics through loop formulas and related constructs, and then extend the language with features such as disjunction and cardinality constraints. These ingredients will be essential when we encode bipolar argumentation in Chapter 5.

4.1 Foundations of Logic Programming

4.1.1 Definite Logic Programs

Before diving into the semantics of Probabilistic Answer Set Programming (PASP), it is important to understand less expressive formalisms, such as *definite* and *propositional* programs. These formalisms underpin Answer Set Programming (ASP): they reveal the inherent complexity of the problems ASP can solve and guide the design of efficient algorithms. We therefore begin by defining the simplest form of logic programming,

namely *definite* (or *positive*) programs:

Definition 4.1.1 (Definite Logic Programs). A definite logic program P is a finite set of clauses (rules) in the form

```
a :- b1, ..., bM.
```

where a, b_1, \dots, b_M are atoms of a function-free FOL L ; and this rule can be seen as material implication restricted to Horn clauses, where $a :- b_1, \dots, b_M$ is read as $B \supset A$ or $B \rightarrow A$ (EITER, IANNI, et al., 2009). The atom a is called the *head* of the rule, while b_1, \dots, b_m are called the rule's *body*. When a rule has an empty body, it is called a fact and can be shortened as a .

The following program is an example of a definite program:

```
happy(turing) :- friends(turing, vonNeumann).
```

Programs without variables, like the one above, are called *propositional* programs.

Although we could omit propositional programs, they are instructive from a complexity perspective: deciding satisfiability for propositional (Horn) programs is P -complete, yet it admits linear-time algorithms (DOWLING and GALLIER, 1984).

4.1.2 Normal Logic Programs

Now, we introduce the concept of programs that include negation in their *body* rules, called Normal:

Definition 4.1.2 (Normal Logic Program). A Normal Logic Program is a finite set of rules of the form:

1 a :- b1, ..., bM, not c1, ..., not cN.

where $M, N \geq 0$; a, b_i, c_i are atoms of a function-free First Order Logic (FOL); and NOT c represents the negation as failure of an atom c ($\neg c$). We also introduce the terminology of *subgoals* in the body, which are either atoms of the form b_i (also called *positive subgoals*) or $\text{not } c_i$ (*negative subgoals*).

4.1.3 Herbrand Universe, Base and Interpretation

To reason about the complexity of logic programs we need additional vocabulary. In particular, the notions of *Herbrand Universe* and *Herbrand Base* allow us to characterise ground instances while balancing expressiveness and efficiency:

Definition 4.1.3 (Herbrand Universe). The *Herbrand Universe* $HU(P)$ of a logic program P is the set of all terms that can be formed from constants and functions in P (with respect to a predefined vocabulary L). The *Herbrand Base* $HB(P)$ of P is the set of all ground atoms that can be formed from terms in $HU(P)$ and predicates of P . Finally, a *Herbrand interpretation* is any subset $I \subseteq HB(P)$ that denotes the ground atoms considered true.

Example 4.1.1. Consider the following logic program P (EITER, IANNI, et al., 2009):

```

h(0, 0).
t(a, b, r).
p(0, 0, b).
p(f(X), Y, Z) :- p(X, Y, Z'), h(X, Y), t(Z, Z', r).
h(f(X), f(Y)) :- p(X, Y, Z'), h(X, Y), t(Z, Z', r).

```

Then, the Herbrand Universe $HU(P)$ is the union of the set containing all constants of P , $\{0, a, b, r\}$, and the set of terms that can be formed from these constants $\{f(0), f(a), f(b), f(r), f(f(0)), f(f(a)), f(f(b)), f(f(r)), \dots\}$. Whereas, the Herbrand Base, $HB(P)$ is given by the set of all ground atoms assertions, $\{p(0, 0, 0), p(a, a, a), \dots, h(0, 0), \dots, t(0, 0, 0), t(a, a, a), \dots\}$.

Finally, we list a few Herbrand Interpretations over $HU(P)$:

- $I_1 = \emptyset$;
- $I_2 = HB(P)$;
- $I_3 = \{h(0, 0), t(a, b, r), p(0, 0, b)\}$.

Note that not all interpretations are consistent with the program P . For instance, I_1 is contradictory because it does not include any of the facts of P .

4.1.4 Grounding of a Logic Program

With the notion of Herbrand Universe and Base, we can now give a more formal definition of *grounding* (EITER, IANNI, et al., 2009):

Definition 4.1.4 (Grounding). A ground instance of a clause C in a logic program P is any clause C' obtained from C by applying a substitution

$$\theta : \text{Var}(C) \rightarrow HU(P),$$

where $\text{Var}(C)$ denotes the variables occurring in C . The grounding of P is the set of all ground instances of its clauses and is denoted by $\text{ground}(P) = \bigcup_{C \in P} \text{ground}(C)$.

4.1.5 Interpretation of a Logic Program

Having established grounding, we next formalise interpretations:

Definition 4.1.5 (Interpretation). The interpretation I of a logic program P is a model of P that is compatible with the assertions in P . That is, I is a model of

- A ground clause $C = a :- b_1, \dots, b_M, \text{not } c_1, \dots, \text{not } c_N$, denoted $I \models C$, if either $\{a, c_1, \dots, c_N\} \cap I \neq \emptyset$ or $\{b_1, \dots, b_M\} \not\subseteq I$;
- A clause C , denoted $I \models C$, if $I \models C'$ for every $C' \in \text{ground}(C)$;
- A program P , denoted $I \models P$, if $I \models C$ for every clause $C \in P$.

That is: an interpretation I is a model of a program P if it is compatible with all the ground instances of the clauses of P .

We call a model I of a program P a *minimal model* if there is no other model J of P such that $J \subset I$. Normal logic programs can have multiple minimal models, whereas definite logic programs admit a unique minimal model (EITER, IANNI, et al., 2009).

4.1.6 Loop Formulas

Definition 4.1.6 (Loops). Let P be a normal logic program. Any nonempty subset L of the atoms of P is called a *loop* if for any two atoms p, q in L , there is a path in the dependency graph of P from p to q of length greater than zero.

We also associate two sets of rules with a loop L :

- The set of rules $R^+(L, P)$ contains rules of P whose heads and bodies are in L ; and
- The set of rules $R_L^-(L, P)$ contains rules about atoms in L that are out of the loop L .

Definition 4.1.7 (Loop Formulas). Let P be a normal logic program and L be a loop of P . Then, the *loop formula* $LF(L, P)$ is the following implication:

$$\bigvee_{p \in L} \neg p \rightarrow \neg \bigwedge_{r \in R_L^-(L, P)} \text{body}(r),$$

where $\text{body}(r)$ is the body of the rule r .

The main idea behind loop formulas is that they can be used to transform logic programs under stable model semantics to propositional theories. This is particularly useful when these propositional theories are fed into SAT solvers to compute stable models. The entire process of translating a logic program by this approach includes: using Clark's completion to create a propositional theory and augment it by using additional loop formulas, that guarantees that this theory admits only stable models (LIN and ZHAO, 2004). One problem with this approach is that the number of loop formulas must be controlled in some way, because there can be an exponential number of loop formulas for a given program (EITER, IANNI, et al., 2009).

4.1.7 Immediate Consequence Operator

To establish the semantics of normal logic programs, we define the notion of the *immediate consequence operator* T_P :

Definition 4.1.8 (T_P Operator). Let P be a Normal Logic Program. The immediate consequence $T_P : 2^{HB(P)} \rightarrow 2^{HB(P)}$ of a set of ground atoms I is defined as:

$$T_P(I) = \left\{ a \mid \begin{array}{l} a :- b_1, \dots, b_M, \text{not } c_1, \dots, c_N \in \text{ground}(P), \\ \text{such that } \{b_1, \dots, b_M\} \subseteq I \text{ and } \{c_1, \dots, c_N\} \cap I = \emptyset \end{array} \right\},$$

where $\text{ground}(P)$ denotes the set of all ground atoms in P .

The most important property of the immediate consequence operator is that it is *monotone* (BOGAERTS and VAN DEN BROECK, 2015) and thus, by the Knaster-Tarki Theorem, has a least fixed point, called $lfp(T_P)$, which is the least model of P (EITER, IANNI, et

al., 2009). Moreover, the sequence obtained by iterating T_P starting from the empty set converges to the fixed point $\text{lfp}(T_P)$.

4.1.8 Negation and Stratification

Even though different classes of Logic Programs, such as Definite and Normal programs, try to capture the intrinsic duality of expressiveness versus efficiency in the field of KR and LPs, these definitions only take into account the different types of rules that can be present in a program, but not their structure.

However, understanding the impact of the negation operator's introduction to program rules is of utmost importance, as it directly creates the possibility of having a collection of multiple minimal models. This change creates the necessity of defining a proper semantic for attributing meaning to negation in logic problems, where there are two main approaches:

- Define a single model for a program, which possibly tries to capture problematic classes of programs due to negation possibly creating multiple minimal models. This approach has success when dealing with *stratified* programs, which we define later. For general normal programs, the most popular semantics following this approach is based on the *well-founded semantics* (that we do not define here for the sake of brevity) (EITER, IANNI, *et al.*, 2009; VAN GELDER *et al.*, 1991).
- Define a collection of models for a program, abandoning the necessity of only one model. This approach gives rise to ASP and its semantics, such as the *stable model semantics*, which we also define later.

One possible way of determining whether a program is *stratified* or not is by finding an ordering for the evaluation of its rules, such that the value of negative literals can be predicted by the values of positive literals (EITER, IANNI, *et al.*, 2009). But we can also define a more general approach by analyzing the structure of the program, that is, the dependency graph:

Definition 4.1.9 (Dependency Graph). Let P be a ground program. Then, its dependency (multi) graph $\text{dep}(P)$ is a tuple of the form (V, E) , where

- A set of nodes V , where each node $v \in V$ is a ground atom occurring in P ; and
- A set of edges E , where each edge $(v, w) \in E$ (v pointing to w , $v \rightarrow w$) occurs if and only if v is the head atom of a rule whose body contains a literal w . If the literal is negative, then the edge is marked as \ast (v, w) ($v \rightarrow^\ast w$).

We call a graph *acyclic* if its grounded dependency graph does not contain any (directed) cycles.

By utilizing the notion of a dependency graph, we can define the concept of stratification:

Definition 4.1.10 (Stratification). Let P be a logic program, and let $G = (V, E)$ be its dependency graph. A stratification of P is a partition $\Sigma = \{S_i \mid i \in \{1, \dots, n\}\}$ of the predicates of P into n non-empty pairwise disjoint sets, such that:

1. If $v \in S_i$, $w \in S_j$, and $(v, w) \in E$, then $i \leq j$ (positive dependency must be within the same stratum or to a higher/later stratum); and
2. If $v \in S_i$ and $*(v, w) \in E$, then $i > j$ (negative dependency must be to a strictly lower/earlier stratum).

A program is called *stratified* if it has some stratification Σ , and its respective partitions S_i are called *strata*.

4.1.9 Stable Model Semantics

As stated before, Normal Logic Programs can have multiple minimal models, which is a consequence of the introduction of negation in the program rules. Following, we have a program where the introduction of negation creates multiple minimal models:

Example 4.1.2. Let P be the following logic program:

```

1  researcher(computability).
2  machine(X) :- researcher(X), not lambda(X).
3  lambda(X) :- researcher(X), not machine(X).

```

Then, note that the program P is not stratified. Furthermore, there is not only one minimal model for P , but two:

1. $M_1 = \{ \text{researcher}(\text{computability}), \text{machine}(\text{computability}) \}$; and
2. $M_2 = \{ \text{researcher}(\text{computability}), \text{lambda}(\text{computability}) \}$.

Hence, to attribute meaning to negation, we first introduce the concept of *reduction* of a program, and then introduce one of the most important semantics for logic programs, the Stable Model semantics.

Definition 4.1.11 (Reduct (short GL-reduct)). Let P be a Normal Logic Program and I be an interpretation. Then, the *reduct* of P w.r.t I , denoted P^I , is obtained by:

1. Grounding the program P ;
2. Removing rules with $\text{not } c$ in the body, for each $c \in I$; and
3. Removing all negative literals $\text{not } c$ from the remaining rules.

Hence, a Gelfond-Lifschitz reduction produces a program that enforces the truth of the atoms in the interpretation I . The first condition, when a literal $c \in I$, makes the negative subgoal $\text{not } c$ false. Consequently, any ground rule containing $\text{not } c$ in its body is removed, as its body cannot be satisfied, and it will not contribute to the least model. When the second condition occurs, if a literal $c \notin I$, it is possible to assume that $\text{not } c$ is indeed true; therefore, it can be removed from the (body of the) rule.

A direct consequence of the definition of the reduct is that the reduction product, P^I , is a positive program. Using this property, it is easy to see that P^I has a least model $LM(P^I)$; and, if P^I is consistent (if it does not present a contradiction), then the least model of P^I is I itself. This bijective relation between P^I and I , of one being able to obtain I from

P^I due to it being a least model of the program, and P^I being the reduct of P w.r.t. I , is the basis of the Stable Model semantics:

Definition 4.1.12 (Stable Model). Let P be a normal logic ground program. An interpretation I is a stable model of P if I is a minimal model of the reduct P^I of P w.r.t. I .

Since a normal program may have several stable models, it is possible to have multiple interpretations that are stable models. Therefore, a normal program may have several stable models (or even none, as we stated earlier). For instance, the example above has two stable models, M_1 and M_2 .

An interesting property of stable models is that they are fixed points of the immediate consequence operator T_P . Formally, a stable model I of P satisfies $T_P(I) = I$ (but the converse is not necessarily true) (EITER, IANNI, et al., 2009).

In the context of this work, we will call stable models *answer sets*. Usually, we call the stable models of a program P *answer sets* when we are talking about Extended Logic Programs (ELP), which can contain *strongly negated* atoms, but we do not define these ELPs, and instead recommend the reader to refer to (EITER, IANNI, et al., 2009).

4.1.10 Reasoning

There are three types of reasoning with logic programs under the stable model semantics. We enumerate them in order of complexity, where the former can be reduced as an instance of the latter:

1. **Consistency** (Satisfiability): decide whether a program has at least one answer set (consistent);
2. **Brave reasoning**: given a ground literal Q , decide whether some (at least one) answer set satisfies Q . If it does, Q is called a *brave consequence* of the program; and
3. **Cautious reasoning**: given a ground literal Q , decide whether all answer sets satisfy Q . If all of them do, Q is called a *cautious consequence*.

4.2 Answer Set Programming

Finally, after defining the main semantics for logic programs that include negation, we introduce the concept of ASP, an extension of normal logic programming that allows for the use of: integrity constraints, strong negation, disjunctive rules, and choice rules (EITER, IANNI, et al., 2009; MAUÁ and COZMAN, 2020).

4.2.1 Extended Logic Programs

We start this ASP definition by first introducing the concept of Extended Logic Programs (ELP) and Extended Disjunctive Logic Programs (EDLP), programs that use the three extensions defined above.

Integrity Constraints

Integrity constraints rule out interpretations that violate hard requirements. They are typically written as headless rules of the form

```
:– b1, ..., bM, not c1, ..., not cN.
```

The same requirement can be expressed with a head by introducing an auxiliary predicate:

```
falsity :- b1, ..., bM, not falsity, not c1, ..., not cN,
```

where the auxiliary *falsity* is a fresh propositional atom (EITER, IANNI, *et al.*, 2009).

Strong Negation

Strong negation is not strictly required—default negation combined with integrity constraints can simulate it—but having a dedicated connective makes programs clearer. While classical logic writes the fact that a is known to be false as $\neg a$, ASP customarily denotes it by $\neg a$.

Extended Logic Programs

By combining normal logic programs with integrity constraints and strong negation, we are capable of defining ELP:

Definition 4.2.1 (Extended Logic Programs). An extended logic program P is a finite set of rules of the form:

```
a :- b1, ..., bM, not c1, ..., not cN,
```

where $M, N \geq 0$; a, b_i, c_i are atoms or *strongly negated* atoms of a FOL.

Because integrity constraints and strong negation can be reduced to the basic semantics, we do not need a new notion of model: stable models of extended programs remain *answer sets*.

Disjunctive Logic Programs

The last extension that we define for normal logic programs is: the addition of disjunctions on the head of rules:

Definition 4.2.2 (Extended Disjunctive Logic Programs). An extended disjunctive logic program P is a finite set of rules of the form:

```
a1 ; ... ; aK :- b1, ..., bM, not c1, ..., not cN,
```

where $K, M, N \geq 0$; a_i, b_i, c_i are atoms or *strongly negated* atoms of a FOL; and the symbol `;` denotes disjunction (similar to how commas denote conjunctions).

The semantics of EDLPs mirror those of extended programs with two additions: answer sets are now minimal models of the reduct P^M (which may have multiple minimal models), and Clark's completion must be generalised to cope with disjunctive heads (ALVIANO, DODARO, *et al.*, 2016).

Finally, we recall the notion of interpretation for EDLPs, unifying the various extensions that lead from normal logic programs to full ASP.

Definition 4.2.3 (Interpretation (of an EDLP)). Let P be an EDLP. An interpretation I is a model of:

1. A ground clause $C = \text{a1}; \dots; \text{aK} :- \text{b1}, \dots, \text{bM}$, $\text{not c1}, \dots, \text{not cN}$, denoted $I \models C$, if either $\{a_1, \dots, a_K, c_1, \dots, c_N\} \cap I \neq \emptyset$ or $\{b_1, \dots, b_M\} \not\subseteq I$;
2. A clause C , denoted $I \models C$, if $I \models C'$ for every $C' \in \text{ground}(C)$; and
3. A program P , denoted $I \models P$, if $I \models C$ for every clause $C \in P$.

4.2.2 Cardinality Constraints

The last extension that characterises Answer Set Programming is the notion of *cardinality constraints*, constructs of the form: $L \backslash \{l_1, \dots, l_n\} \cup$, that are satisfied whenever the number of satisfied literals l_i is between the integral bounds L and U , inclusive (SYRJÄNEN and NIEMELÄ, 2001). As Syrjänen and Niemelä described, a cardinality constraint in a rule head imposes a non-deterministic choice over the literals in it when the rule body is satisfied.

Choice Rules

A special case of cardinality constraints are *choice rules*: rules where the head is enclosed in brackets and represent the idea that the head can be included in a stable model only if the body holds; but it can be left out, too (SYRJÄNEN and NIEMELÄ, 2001). This type of rule can be expressed through normal rules by introducing a new atom. For example,

Example 4.2.1. The following program P

$\{a\} :- b, \text{not } c.$

is equivalent to the Normal logic program P' :

$a :- \text{not aa}, b, \text{not } c.$
 $\text{aa} :- \text{not } a.$

4.3 Probabilistic Logic Programming

To define Probabilistic Answer Set Programming, we first review how logic programs can be enriched to express probability distributions as generative models. Sato's distribution semantics (SATO, 1995) is one of the seminal works in this area: it associates probability to a set of *independent* events in such a way that a unique probability measure is induced over all interpretations of ground atoms (COZMAN and MAUÁ, 2020).

Definition 4.3.1 (Probabilistic Fact). A probabilistic fact is a pair consisting of an atom A and a probability value α , which we denote by $\alpha :: A$ (COZMAN and MAUÁ, 2017).

Additionally, atoms of probabilistic facts may contain logical variables ($\alpha :: r(X_1, \dots, X_n)$), and are therefore ungrounded. When considering this case, we interpret

such a probabilistic fact as the set of all grounded probabilistic facts, obtained after grounding the variables in the atom.

Finally, probabilistic facts do not unify with the head of any rule: there is no substitution θ that makes a probabilistic fact f and a rule head h syntactically identical, i.e. $f\theta \neq h\theta$ for all choices of f and h .

Note that this definition of PLPs only associates probabilities through the use of probabilistic facts. Thus, properties from logic programs are also present in PLPs, such as *stratification*, *acyclicity*, and *definiteness*.

In order to further extend this probabilistic semantics, one may define multi-valued scenarios, where probabilities are not of the type $p::a$, where a has probability p of being *true* and $(1-p)$ of being *false*; but instead we have a distribution over many scenarios $a(1), \dots, a(k)$, each associated with a probability $p(i)$. This notion, called *annotated disjunctions*, extends the previous definition of PLPs to Extended Disjunctive Logic Programming (EDLP) by adding probabilistic facts to the head of rules.

Definition 4.3.2 (Annotated Disjunctive Rules). A probabilistic extension of an EDLP is called Annotated Disjunctive Rules (ADR), a set of rules of the form:

$$p_1::a(1); \dots; p_k::a(K) :- b_1, \dots, b_M, \text{not } c_1, \dots, \text{not } c_N,$$

where p_1, \dots, p_k are nonnegative real values whose sum is smaller than or equal to 1. In this sense, a probabilistic fact of the form

$$p::a$$

is a special case (and syntactic sugar) of an ADR, where $p::a$ is equivalent to the rule $p::a; (1-p)::\text{dummy} :- ..$, with a fresh atom dummy that is never derived (GEH *et al.*, 2024a).

4.3.1 Probabilistic Answer Set Programming

Although we have briefly defined Sato's distribution semantics, a more formal description is required here, together with adaptations that make the semantics compatible with ASP programs under stable models.

It is important to emphasise that probabilities are assigned to entire PLPs, not to individual logical consequences. This makes probabilistic semantics over PLPs agnostic with respect to the underlying logical semantics, which is particularly useful when considering non-stratified programs that may use a *stable model* or *well-founded* semantics (EITER, IANNI, *et al.*, 2009; MAUÁ and COZMAN, 2020; GEH *et al.*, 2024b).

Definition 4.3.3 (Probabilistic Choice). A PLP (P, PF) with n probabilistic facts can generate a total of 2^n different logic programs: one for each possible assignment of the probabilistic facts (when considering binary probabilistic facts). The assignment of a probabilistic fact $\alpha :: A$ can be seen as choosing to keep or erase the fact A from the program (COZMAN and MAUÁ, 2017).

As commented before, these *probabilistic choices* are assumed to be independent.

Definition 4.3.4 (Total Choice). A *total choice* θ for a PLP (P, PF) is a subset of the set of grounded probabilistic facts. Thus, θ can be interpreted as a set of ground facts that are probabilistically selected to be included in P , while all other ground facts obtained from probabilistic facts are discarded (COZMAN and MAUÁ, 2017).

Because the probabilistic choices are assumed to be independent, the probability of a total choice can be easily computed by the product over the grounded probabilistic facts, where a probabilistic fact $\alpha :: A$ that is chosen to be included contributes with a factor of α , and, if it is discarded, it contributes with a factor of $(1 - \alpha)$ (COZMAN and MAUÁ, 2017).

An important property of total choices is that, for a total choice θ , the induced logic program (w.r.t. P), denoted by $P \cap PF^{\downarrow\theta}$, is a normal logic program (COZMAN and MAUÁ, 2017; GEH *et al.*, 2024a). Therefore, for an ADR, a total choice θ induces a normal logic program, where each rule r from the ADR is transformed into a rule of the form

```
a_i :- b1, ..., bM, not c1, ..., not cN,
```

where i is the corresponding index of the probabilistic choice in θ (GEH *et al.*, 2024a).

Whenever we have a total choice θ for a PLP such that $P \cup PF^{\downarrow\theta}$ is stratified, the resulting induced program has only one *stable model*, as we can conclude by earlier subsections, particularly 4.1.9. Hence, Sato's distribution semantics can be naturally extended to this class of programs, and does not create any controversial interpretation of the probabilistic semantics.

A more interesting case is when we have a total choice θ for a PLP that induces a (non-stratified) logical program with multiple *stable models*. In this case, we have the need to define semantics capable of handling this situation: the Credal and MAXENT semantics, which we define next.

First, we define the notion of *consistency* for PLPs:

Definition 4.3.5 (Consistency of Probabilistic Logic Programs). A PLP (P, PF) is *consistent* if, for each total choice θ , the induced logic program $P \cup PF^{\downarrow\theta}$ has at least one stable model (COZMAN and MAUÁ, 2017).

Given that we have a definition to express PLPs that are able to have stable models, independent of the total choice, we are capable of formalizing the notion of a *probability model* for PLPs:

Definition 4.3.6 (Probability Model for Probabilistic Logic Programs). A *probability model* for a consistent PLP (P, PF) is a probability measure \mathbb{P} over interpretations of P , such that (COZMAN and MAUÁ, 2017):

1. Every interpretation I with $\mathbb{P}(I) > 0$ is a stable model of (the normal logic program $P \cup PF^{\downarrow\theta}$ w.r.t. the total choice θ that agrees with I on the probabilistic facts; and
2. The probability of each total choice θ is the product of the probabilities for all individual choices in θ :

$$\mathbb{P}(\{ I \mid I \cap C = C \}) = \prod_{\alpha :: A \mid A \in C} \alpha \prod_{\alpha :: A \mid A \notin C} (1 - \alpha).$$

Credal Semantics

As the first PASP semantics to be described, we introduce the *credal semantics*, which is based on the idea of using intervals (containing *upper* and *lower*) to represent the uncertainty of the probability (LUKASIEWICZ, 2005; LUKASIEWICZ, 2007).

Definition 4.3.7 (Credal Semantics). The *credal semantics* for a consistent PLP (P, PF) is the set of all probability models for (P, PF) (COZMAN and MAUÁ, 2017).

Moreover, the credal semantics is a *closed* and *convex* set of probability measures, and corresponds to the set of all probability measures that dominate an infinitely monotone Choquet capacity (COZMAN and MAUÁ, 2017) (more about Choquet capacities can be found in COZMAN and MAUÁ, 2020). As such, the credal semantics of any set of models \mathcal{M} is characterized by an interval $[\underline{\mathbb{P}}(\mathcal{M}), \bar{\mathbb{P}}(\mathcal{M})]$:

$$\underline{\mathbb{P}}(\mathcal{M}) = \sum_{\theta \in \Theta : \Gamma(\theta) \subseteq \mathcal{M}} \mathbb{P}(\theta), \quad \bar{\mathbb{P}}(\mathcal{M}) = \sum_{\theta \in \Theta : \Gamma(\theta) \cap \mathcal{M} \neq \emptyset} \mathbb{P}(\theta), \quad (4.1)$$

where Θ is the set of all total choices, and $\Gamma(\theta)$ is the set of stable models associated with the total choice (COZMAN and MAUÁ, 2017; COZMAN and MAUÁ, 2020).

Given a PLP, this interval can be computed for a set of assignments \mathbf{Q} for ground atoms by the following algorithm (COZMAN and MAUÁ, 2017):

1. Given a PLP (P, PF) and \mathbf{Q} , initialize variables a and b with 0;
2. For each total choice θ , compute the set S of all stable models of $P \cap PF^{\downarrow\theta}$, and:
 - (a) If \mathbf{Q} is *true* in every stable model in S , then $a \leftarrow a + P(\theta)$; and
 - (b) If \mathbf{Q} is *true* in some stable model of S , then $b \leftarrow b + P(\theta)$.
3. Return $[a, b]$ as the interval $[\underline{\mathbb{P}}(\mathbf{Q}), \bar{\mathbb{P}}(\mathbf{Q})]$.

An interesting relation between the credal semantics and reasoning over Answer Sets is that the computation of the upper probability, $\bar{\mathbb{P}}(\mathbf{Q})$, involves *brave reasoning*, and the computation of the lower probability, $\underline{\mathbb{P}}(\mathbf{Q})$, involves *cautious reasoning* (COZMAN and MAUÁ, 2017).

Maximum-Entropy Semantics

To conclude this chapter, we define one of the most famous semantics for PASP: the *Maximum Entropy* semantics, also known as MAXENT semantics:

Definition 4.3.8 (Maximum Entropy (MaxEnt) Semantics). This semantics is based on the principle of maximum entropy: evenly dividing the probability mass among all stable models of a program (induced by a total choice). Formally, for a stable model \mathcal{M} , the probability $\mathbb{P}(\mathcal{M})$ is given by:

$$\mathbb{P}(\mathcal{M}) = \sum_{\theta : \mathcal{M} \in \Gamma(\theta)} \frac{\mathbb{P}(\theta)}{n},$$

where n is the number of stable models associated with the total choice θ (Cozman and Mauá, 2017).

4.3.2 PASP Inference as 2AMC

By analyzing the definitions of PASP semantics, an immediate observation is that they are closely related to the *two-level weighted model counting* (2AMC) problem defined in Chapter 3. In fact, both the credal semantics and the MAXENT semantics can be reduced to instances of 2AMC (Kiesel and Eiter, 2023; Azzolini and Riguzzi, 2023). In more detail, when defining the 2AMC framework, (Kiesel and Eiter, 2023) showed that the MAXENT semantics can be directly mapped to a 2AMC instance by having an inner model counting problem that counts the number of stable models for each total choice, and an outer problem that sums the probabilities of these total choices. Similarly, (Azzolini and Riguzzi, 2023) demonstrated that the credal semantics can also be expressed as a 2AMC problem, where the inner problem applies *brave/cautious* reasoning to verify whether the query holds in some/all stable models for each total choice, and the outer problem sums the probabilities of these total choices, as a “standard” weighted model counting problem.

Chapter 5

Contribution

We pursue a neuro-symbolic pipeline for argumentation mining that leverages probabilistic Answer Set Programming (PASP) as an expressive representation language and compiles PASP programs into tractable circuits for efficient inference. By embedding argumentative constraints into compiled circuits we aim to align neural predictions with the semantics of bipolar argumentation, enabling transparent reasoning, calibrated uncertainty, and scalable learning loops.

5.1 Scope and Roadmap

This chapter is organised in four stages. Section 5.2 formalises the problem we study, first by explaining how textual argumentation signals are mapped into PASP programs and then by detailing the knowledge compilation pipeline that supports tractable inference. Section 5.3 gathers the foundational notions we rely on: we introduce the argumentation structures we model, recall Clark’s completion, and discuss the circuit properties we require from downstream compilers. The chapter closes with Section 5.4, which anticipates the experimental plan that will ultimately assess the proposed approach.

5.2 Problem Setting

Before delving into the technical building blocks of our approach, we first outline the problem setting we target and summarise the proposed pipeline for scalable neuro-symbolic argumentation mining.

5.2.1 From Argumentation Mining to PASP

Argumentation mining systems must recognise argumentative spans, classify their roles, and recover the interaction graph that ties them together (STAB and GUREVYCH, 2017). Early neuro-symbolic pipelines coupled local neural predictions with Integer Linear Programming constraints, whereas probabilistic logic programming brought principled uncertainty handling to the task (FIERENS *et al.*, 2015; MANHAEVE *et al.*, 2018; CERVEIRA

DO AMARAL *et al.*, 2023). Stable model semantics further broadens the expressiveness of these pipelines by accommodating negative cycles and other non-monotonic phenomena that arise in real discourse (TOTIS *et al.*, 2023).

We formalise an argumentation mining instance as a bipolar argumentation framework whose nodes represent candidate arguments extracted from text and whose labelled edges denote support or attack relations (TONI and SERGOT, 2011; TONI, POTYKA, *et al.*, 2023). Neural components provide noisy observations—e.g. span detection probabilities or relational scores—that we encode as probabilistic facts. The deterministic backbone of the PASP program asserts domain constraints, propagates support chains, and enforces well-formed argumentative structures. Inference tasks such as cautious or brave reasoning over accepted arguments then reduce to answering PASP queries, which we cast as instances of second-level algebraic model counting (2AMC) (KIESEL, TOTIS, *et al.*, 2022).

5.2.2 Knowledge Compilation Pipeline

Directly solving these probabilistic programs on demand would be prohibitively expensive. Instead, we adopt a knowledge compilation workflow that proceeds in three stages: grounding, propositional rewriting, and circuit construction. Grounding produces a finite propositional theory while retaining the structure of the original program; propositional rewriting augments the theory with loop formulas and other constraints that guarantee equivalence under stable-model semantics; and circuit construction compiles the resulting theory into a structured representation that supports efficient 2AMC evaluation.

The compiled artefact supports repeated 2AMC evaluations: evidence and query updates modify only the literal labels, after which the circuit can be re-evaluated in time linear in its size (KIESEL, TOTIS, *et al.*, 2022).

5.3 Foundational Building Blocks

In this section, we review the foundational concepts that underpin our proposed pipeline, starting with bipolar argumentation frameworks for modelling argumentative structures in PASP, and then discussing the knowledge compilation techniques that enable efficient inference in this setting.

5.3.1 Argumentation Frameworks

Bipolar argumentation frameworks extend Dung’s abstract argumentation with explicit support relations that can interact with attacks in non-trivial ways.

Definition 5.3.1 (Bipolar Argumentation Framework). A bipolar argumentation framework (BAF) is a triple $\langle A, R_d, R_s \rangle$ where A is a set of arguments, $R_d \subseteq A \times A$ is the defeat (attack) relation, and $R_s \subseteq A \times A$ is the support relation. Support chains can propagate attacks: a path of supports followed by a defeat induces a derived attack along the support path (TONI and SERGOT, 2011; TONI, POTYKA, *et al.*, 2023).

Encoding a BAF in PASP lets us reason about accepted arguments under probabilistic stable-model semantics while accounting for uncertainty in the edges or their textual

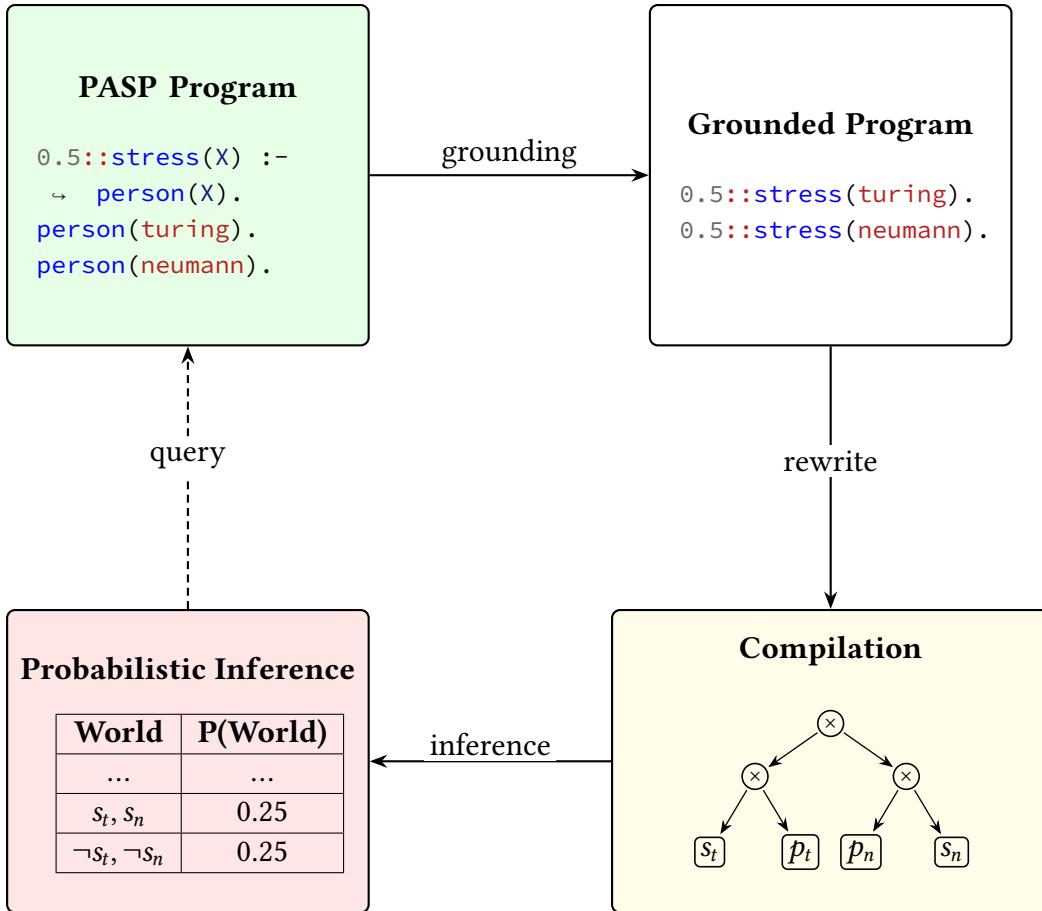


Figure 5.1: Knowledge compilation pipeline with grounding, compilation, and inference stages.

evidence.

5.3.2 Knowledge Compilation

Once the PASP program is grounded, we apply Clark's completion and generate a structured circuit that satisfies the properties required for tractable 2AMC inference.

Clark's Completion and CNF Generation

In order to obtain a propositional theory equivalent to the stable-model semantics of the grounded program, we first apply Clark's completion (CLARK, 1978). For normal programs, the Clark's completion equates each atom with the disjunction of the bodies of the rules that derive it. For disjunctive programs we obtain

$$\text{CLARK}(P) = \bigwedge_{a \in \mathcal{A}(P)} \left[a \iff \bigvee_{r \in \mathcal{R}(P,a)} \bigwedge_{b \in \text{body}(r)} b \bigwedge_{a' \in \text{head}(r) \setminus \{a\}} \neg a' \right], \quad (5.1)$$

which yields a conjunctive normal form (CNF) after distributing conjunctions and disjunctions over the finite grounding. Loop formulas supplement the completion to eliminate unsupported models, guaranteeing equivalence with the stable-model semantics (LIN

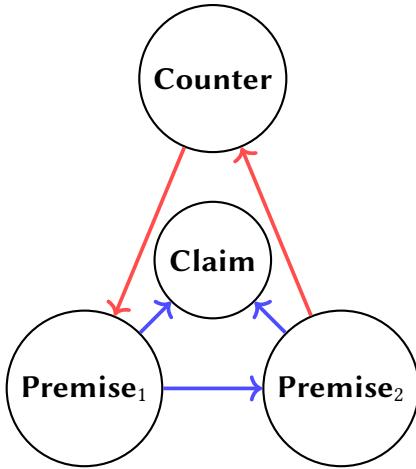


Figure 5.2: Illustrative bipolar argumentation fragment with mutual support and derived defeat chains.

and ZHAO, 2004; EITER, IANNI, *et al.*, 2009). Recent analyses tighten this translation by bounding the number of auxiliary variables introduced during completion, which is crucial for predictable compilation cost (EITER, HECHER, *et al.*, 2024).

Circuits

For the purposes of this work, Logic circuits are directed acyclic graphs that represent Boolean functions through internal AND and OR nodes, with literals labelling the leaves. This family of circuits are often called Negation Normal Form (NNF) circuits (Adnan DARWICHE and Pierre MARQUIS, 2002), and there are two key approaches to using them for representing complex boolean formulas: top-down compilation, which first translates the formula into a specific normal form (e.g., CNF) and then compiles it into a circuit; and bottom-up compilation, which constructs the circuit directly from the formula using a set of compilation rules (by applying operations such as conjunction, disjunction, and negation).

To support efficient 2AMC inference, we require circuits that satisfy three main properties that make algebraic model counting tractable (A. DARWICHE and P. MARQUIS, 2002; Adnan DARWICHE, 2011):

- **Decomposability:** sub-circuits combined by a product node refer to disjoint sets of variables.
- **Determinism:** the children of a sum node are mutually exclusive, preventing double counting.
- **Smoothness:** the children of a sum node mention the same variables, enabling weight sharing and differentiation.

Circuits satisfying these properties belong to the sd-DNNF class (smooth, deterministic, decomposable negation normal form). State-of-the-art compilers such as C2D, D4, and SHARPSAT-TD can target this class while preserving the structural guarantees required for efficient 2AMC inference (EITER, HECHER, *et al.*, 2024). Their output supports on-the-fly evaluation of credal and MAXENT semantics through algebraic model counting (WANG *et al.*,

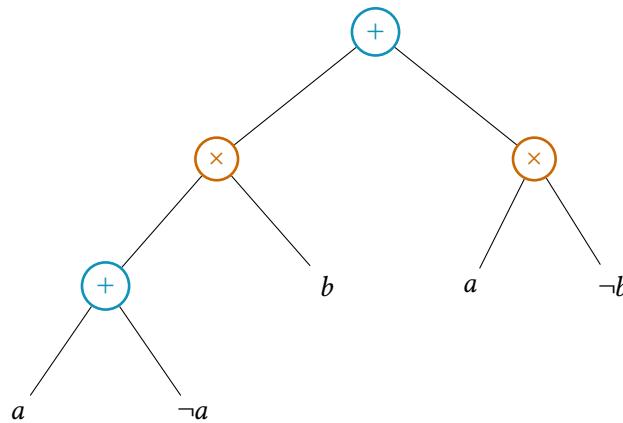


Figure 5.3: Smooth, decomposable, and deterministic circuit fragment used for Pasp inference.

2025). Moreover, the recent translation scheme of MAENE *et al.* (2024) converts logic circuits into tensor-based representations compatible with automatic differentiation frameworks such as PyTorch or JAX, enabling gradient-based learning over the compiled model.

5.4 Experimental Results

The experimental section of the dissertation will benchmark the proposed pipeline against established sd-DNNF compilers. We plan to evaluate C2D, D4, and SHARPSAT-TD, three top-down compilers adapted by (EITER, HECHER, *et al.*, 2024) in order to handle the necessary constraints for tractable Pasp inference as 2AMC problem (WANG *et al.*, 2025). The workloads will include the an argumentation program family, which captures the argumentation instances that motivated our pipeline.

First, we describe the experimental pipeline that will be used to evaluate the selected knowledge compilers on argumentation mining instances. This outlook clarifies that the compiled circuits will be assessed mainly for compilation efficiency, as compiled circuits can be reused for multiple inference tasks during neuro-symbolic learning. Then, we discuss the expected outcomes of the experiments, focusing on the research question that the evaluation aims to address. Finally, we analyse the experimental results obtained from running the proposed pipeline with the selected compilers in argumenatation mining instances ranging from 3 nodes to 50 nodes (a large scale when considering the complexity of exact neuro-symbolic inference).

5.4.1 Experimental Outlook

Program listing. The generator produces instantiations of the following ungrounded Pasp template, which encodes argumentation constraints before grounding:

```
% Domain and probabilistic choices over components
node(ARG) :- argument(ARG).
0.5::aux_component(ARG,0) :- node(ARG).
0.5::aux_component(ARG,1) :- node(ARG).
```

```

component(ARG,0) :- aux_component(ARG,0).
component(ARG,1) :- aux_component(ARG,1), not aux_component(ARG,0).
component(ARG,2) :- node(ARG), not aux_component(ARG,1), not
    aux_component(ARG,0).

% Probabilistic relations between distinct arguments
0.5::relation(SRC,DST) :- node(SRC), node(DST), SRC != DST.

% Stratification constraints
fail :- node(SRC), node(DST), SRC != DST, component(SRC,2), not
    relation(SRC,DST).
fail :- node(SRC), node(DST), SRC != DST, component(SRC,1), not
    component(DST,2), not relation(SRC,DST).
fail :- node(SRC), node(DST), SRC != DST, component(SRC,0), component(DST,2),
    not relation(SRC,DST).

```

Grounding. The grounding step instantiates the above program over a set of arguments provided as input. For an input size of n arguments, the grounded program contains $3n$ probabilistic facts for the components and $n(n - 1)/2$ probabilistic facts for the relations, resulting in a total of $3n + n(n - 1)/2$ probabilistic facts. Note that the number of rules in the grounded program grows quadratically with n due to the pairwise relations between arguments. In order to compute such groundings, we use *clingo*, the state-of-the-art ASP grounder ([GEBSER et al., 2014](#)).

Knowledge compilation. Given the grounded program, we apply Clark’s completion and generate the corresponding CNF with loop formulas, by following the algorithm present in ([LEE and LIFSHITZ, 2003](#)) (another possible approach would be to follow the algorithm present in ([EITER, HECHER, et al., 2024](#)), which uses the cycle-breaking method from ([EITER, HECHER, et al., 2021](#))). The resulting CNF is then fed to the selected knowledge compilers, which produce sd-DNNF circuits that support efficient 2AMC inference.

Evaluation metrics. We will assess the performance of the knowledge compilers along two axes: compilation time and circuit size. Compilation time measures the duration required by each compiler to process the CNF and produce the sd-DNNF circuit. Circuit size is evaluated in terms of the number of nodes and edges in the resulting circuit, as well as the memory footprint. These metrics provide insights into the efficiency and scalability of each compiler when handling the argumentation instances generated by our program template.

Neuro-Symbolic Learning Given that the circuits produced by the compilers support efficient 2AMC inference, they do not need to be recompiled when evidence or queries change, i.e. then the probabilities estimated by neural components vary during training. This property enables the integration of the compiled circuits into end-to-end neuro-symbolic learning loops, where neural networks provide probabilistic inputs to the PASP program, and gradients can be backpropagated through the circuit to update the neural parameters. We plan to explore this integration in future work, building upon the tensor-based circuit representations introduced by [MAENE et al. \(2024\)](#).

5.4.2 Expected Outcomes

The only variable factor in the proposed pipeline is the knowledge compiler, as the grounding and CNF generation steps are deterministic and the PASP program is fixed. Therefore, we expect the experimental evaluation to reveal differences in compilation time and circuit size among the selected compilers. These differences will inform the choice of the most suitable compiler for integrating into neuro-symbolic learning loops for argumentation mining. Ultimately, the goal is to identify a compilation strategy that balances efficiency and scalability, enabling practical applications of probabilistic argumentation mining in real-world scenarios.

Therefore, the only applicable research question for this experimental outlook is: which knowledge compiler among C2D, D4, and SHARPSAT-TD produces the most efficient sd-DNNF circuits for argumentation instances, in terms of compilation time and circuit size?

Given the results from ([EITER, HECHER, et al., 2024](#)) on a variety of PASP programs, we expect SHARPSAT-TD to outperform the other two compilers in both compilation time and circuit size, due to its advanced top-down compilation techniques, while C2D and D4 may exhibit longer compilation times and larger circuits. However, the actual performance may vary depending on the specific structure of the argumentation instances generated by our program template.

5.4.3 Empirical Analysis

References

- [AKERS 1978] AKERS. “Binary decision diagrams”. *IEEE Transactions on computers* 100.6 (1978), pp. 509–516 (cit. on p. 3).
- [ALVIANO, DODARO, *et al.* 2016] Mario ALVIANO, Carmine DODARO, *et al.* “Completion of disjunctive logic programs.” In: *IJCAI*. Vol. 16. 2016, pp. 886–892 (cit. on p. 20).
- [ARORA and BARAK 2009] Sanjeev ARORA and Boaz BARAK. *Computational complexity: a modern approach*. Cambridge University Press, 2009 (cit. on pp. 3–5).
- [AZZOLINI and RIGUZZI 2023] Damiano AZZOLINI and Fabrizio RIGUZZI. “Inference in probabilistic answer set programming under the credal semantics”. In: *International Conference of the Italian Association for Artificial Intelligence*. Springer. 2023, pp. 367–380 (cit. on p. 25).
- [BOGAERTS and VAN DEN BROECK 2015] Bart BOGAERTS and Guy VAN DEN BROECK. “Knowledge compilation of logic programs using approximation fixpoint theory”. *Theory and Practice of Logic Programming* 15.4-5 (2015), pp. 464–480 (cit. on p. 16).
- [BRYANT 1986] Randal E BRYANT. “Graph-based algorithms for boolean function manipulation”. *Computers, IEEE Transactions on* 100.8 (1986), pp. 677–691 (cit. on p. 3).
- [CERVEIRA DO AMARAL *et al.* 2023] Filipe CERVEIRA DO AMARAL, H Sofia PINTO, and Bruno MARTINS. “Argumentation mining from textual documents combining deep learning and reasoning”. In: *EPIA Conference on Artificial Intelligence*. Springer. 2023, pp. 389–401 (cit. on pp. 1, 27).
- [CHAKRABORTY *et al.* 2015] Supratik CHAKRABORTY, Dror FRIED, Kuldeep S MEEL, and Moshe Y VARDI. “From weighted to unweighted model counting.” In: *IJCAI*. 2015, pp. 689–695 (cit. on p. 6).
- [CHAVIRA and Adnan DARWICHE 2008] Mark CHAVIRA and Adnan DARWICHE. “On probabilistic inference by weighted model counting”. *Artificial Intelligence* 172.6-7 (2008), pp. 772–799 (cit. on pp. 7, 8).

- [CHAVIRA, Adnan DARWICHE, and JAEGER 2006] Mark CHAVIRA, Adnan DARWICHE, and Manfred JAEGER. “Compiling relational bayesian networks for exact inference”. *International Journal of Approximate Reasoning* 42.1-2 (2006), pp. 4–20 (cit. on pp. 7, 8).
- [CLARK 1978] Keith L CLARK. “Negation as failure”. In: *Logic and Databases*. Ed. by Hervé GALLAIRE and Jack MINKER. Springer, 1978, pp. 293–322 (cit. on p. 29).
- [COOK 1971] Stephen A. COOK. “The complexity of theorem-proving procedures”. In: *Proceedings of the Third Annual ACM Symposium on Theory of Computing*. STOC ’71. New York, NY, USA: Association for Computing Machinery, 1971, pp. 151–158. ISBN: 9781450374644. DOI: [10.1145/800157.805047](https://doi.org/10.1145/800157.805047). URL: <https://doi.org/10.1145/800157.805047> (cit. on pp. 4, 6).
- [COZMAN and MAUÁ 2017] Fabio Gagliardi COZMAN and Denis Deratani MAUÁ. “On the semantics and complexity of probabilistic logic programs”. *Journal of Artificial Intelligence Research* 60 (2017), pp. 221–262 (cit. on pp. 21–25).
- [COZMAN and MAUÁ 2020] Fabio Gagliardi COZMAN and Denis Deratani MAUÁ. “The joy of probabilistic answer set programming: semantics, complexity, expressivity, inference”. *International Journal of Approximate Reasoning* 125 (2020), pp. 218–239 (cit. on pp. 21, 24).
- [A. DARWICHE and P. MARQUIS 2002] A. DARWICHE and P. MARQUIS. “A knowledge compilation map”. *Journal of Artificial Intelligence Research* 17 (Sept. 2002), pp. 229–264. ISSN: 1076-9757. DOI: [10.1613/jair.989](http://dx.doi.org/10.1613/jair.989). URL: <http://dx.doi.org/10.1613/jair.989> (cit. on p. 30).
- [Adnan DARWICHE 2002] Adnan DARWICHE. “A logical approach to factoring belief networks”. *KR* 2 (2002), pp. 409–420 (cit. on pp. 7, 8).
- [Adnan DARWICHE 2011] Adnan DARWICHE. “Sdd: a new canonical representation of propositional knowledge bases”. In: *Twenty-Second International Joint Conference on Artificial Intelligence*. 2011 (cit. on p. 30).
- [Adnan DARWICHE and Pierre MARQUIS 2002] Adnan DARWICHE and Pierre MARQUIS. “A knowledge compilation map”. *Journal of Artificial Intelligence Research* 17 (2002), pp. 229–264 (cit. on p. 30).
- [DOWLING and GALLIER 1984] William F DOWLING and Jean H GALLIER. “Linear-time algorithms for testing the satisfiability of propositional horn formulae”. *The Journal of Logic Programming* 1.3 (1984), pp. 267–284 (cit. on p. 14).
- [EITER, HECHER, et al. 2021] Thomas EITER, Markus HECHER, and Rafael KIESEL. “Treewidth-aware cycle breaking for algebraic answer set counting”. In: *Proceedings of the International Conference on Principles of Knowledge Representation and Reasoning*. Vol. 18. 1. 2021, pp. 269–279 (cit. on p. 32).

REFERENCES

- [EITER, HECHER, *et al.* 2024] Thomas EITER, Markus HECHER, and Rafael KIESEL. “Aspmc: new frontiers of algebraic answer set counting”. *Artificial Intelligence* 330 (2024), p. 104109. ISSN: 0004-3702. DOI: <https://doi.org/10.1016/j.artint.2024.104109>. URL: <https://www.sciencedirect.com/science/Article/pii/S0004370224000456> (cit. on pp. 30–33).
- [EITER, IANNI, *et al.* 2009] Thomas EITER, Giovambattista IANNI, and Thomas KRENNWALLNER. *Answer set programming: A primer*. Springer, 2009 (cit. on pp. 14–17, 19, 20, 22, 30).
- [FIERENS *et al.* 2015] Daan FIERENS *et al.* “Inference and learning in probabilistic logic programs using weighted boolean formulas”. *Theory and Practice of Logic Programming* 15.3 (2015), pp. 358–401 (cit. on pp. 1, 27).
- [FRANCO and MARTIN 2009] John FRANCO and John MARTIN. “A history of satisfiability.” *HandBook of satisfiability* 185 (2009), pp. 3–74 (cit. on p. 3).
- [GEBSER *et al.* 2014] Martin GEBSER, Roland KAMINSKI, Benjamin KAUFMANN, and Torsten SCHAUB. “Clingo= asp+ control: preliminary report”. *arXiv preprint arXiv:1405.3694* (2014) (cit. on p. 32).
- [GEH *et al.* 2024a] Renato Lui GEH, Jonas GONÇALVES, Igor C SILVEIRA, Denis D MAUÁ, and Fabio G COZMAN. “Dpasp: a probabilistic logic programming environment for neurosymbolic learning and reasoning”. In: *Proceedings of the International Conference on Principles of Knowledge Representation and Reasoning*. Vol. 21. 1. 2024, pp. 731–742 (cit. on pp. 22, 23).
- [GEH *et al.* 2024b] Renato Lui GEH, Jonas GONÇALVES, Igor C SILVEIRA, Denis D MAUÁ, and Fabio G COZMAN. “Dpasp: a probabilistic logic programming environment for neurosymbolic learning and reasoning”. In: *Proceedings of the International Conference on Principles of Knowledge Representation and Reasoning*. Vol. 21. 1. 2024, pp. 731–742 (cit. on p. 22).
- [KIESEL and EITER 2023] Rafael KIESEL and Thomas EITER. “Knowledge compilation and more with sharpsat-td”. In: *Proceedings of the 20th International Conference on Principles of Knowledge Representation and Reasoning*. IJCAI Organization, 2023, pp. 406–416 (cit. on p. 25).
- [KIESEL, TOTIS, *et al.* 2022] Rafael KIESEL, Pietro TOTIS, and Angelika KIMMIG. “Efficient knowledge compilation beyond weighted model counting”. *Theory and Practice of Logic Programming* 22.4 (2022), pp. 505–522 (cit. on pp. 7, 10, 28).
- [KIMMIG *et al.* 2017] Angelika KIMMIG, Guy VAN DEN BROECK, and Luc DE RAEDT. “Algebraic model counting”. *Journal of Applied Logic* 22 (2017), pp. 46–62 (cit. on pp. 3, 7, 9).

- [KRENTEL 1988] Mark W. KRENTEL. “The complexity of optimization problems”. *Journal of Computer and System Sciences* 36.3 (1988), pp. 490–509. ISSN: 0022-0000. DOI: [https://doi.org/10.1016/0022-0000\(88\)90039-6](https://doi.org/10.1016/0022-0000(88)90039-6). URL: <https://www.sciencedirect.com/science/Article/pii/0022000088900396> (cit. on p. 5).
- [LEE and LIFSCHITZ 2003] Joohyung LEE and Vladimir LIFSCHITZ. “Loop formulas for disjunctive logic programs”. In: *Logic Programming: 19th International Conference, ICLP 2003, Mumbai, India, December 9-13, 2003. Proceedings 19*. Springer. 2003, pp. 451–465 (cit. on p. 32).
- [LEVESQUE 1986] Hector J LEVESQUE. “Knowledge representation and reasoning”. *Annual review of computer science* 1.1 (1986), pp. 255–287 (cit. on p. 13).
- [LEVIN 1973] Leonid Anatolevich LEVIN. “Universal sequential search problems”. *Problemy peredachi informatsii* 9.3 (1973), pp. 115–116 (cit. on pp. 4, 6).
- [LIN and ZHAO 2004] Fangzhen LIN and Yuting ZHAO. “Assat: computing answer sets of a logic program by sat solvers”. *Artificial Intelligence* 157.1-2 (2004), pp. 115–137 (cit. on pp. 16, 29).
- [LUKASIEWICZ 2005] Thomas LUKASIEWICZ. “Probabilistic description logic programs”. In: *European Conference on Symbolic and Quantitative Approaches to Reasoning and Uncertainty*. Springer. 2005, pp. 737–749 (cit. on p. 24).
- [LUKASIEWICZ 2007] Thomas LUKASIEWICZ. “Probabilistic description logic programs”. *International Journal of Approximate Reasoning* 45.2 (2007), pp. 288–307 (cit. on p. 24).
- [MAENE *et al.* 2024] Jaron MAENE, Vincent DERKINDEREN, and Pedro Zuidberg Dos MARTIRES. “Klay: accelerating neurosymbolic ai”. *arXiv preprint arXiv:2410.11415* (2024) (cit. on pp. 31, 32).
- [MANHAEVE *et al.* 2018] Robin MANHAEVE, Sebastijan DUMANCIC, Angelika KIMMIG, Thomas DEMEESTER, and Luc DE RAEDT. “Deepproblog: neural probabilistic logic programming”. *Advances in neural information processing systems* 31 (2018) (cit. on pp. 1, 27).
- [MAUÁ and COZMAN 2020] Denis Deratani MAUÁ and Fabio Gagliardi COZMAN. “Complexity results for probabilistic answer set programming”. *International Journal of Approximate Reasoning* 118 (2020), pp. 133–154 (cit. on pp. 19, 22).
- [SANG *et al.* 2005] Tian SANG, Paul BEAME, and Henry A KAUTZ. “Performing bayesian inference by weighted model counting”. In: *AAAI*. Vol. 5. 2005, pp. 475–481 (cit. on p. 8).
- [SATO 1995] Taisuke SATO. “A statistical learning method for logic programs with distribution semantics” (1995) (cit. on p. 21).

REFERENCES

- [SHANNON 1938] Claude E SHANNON. “A symbolic analysis of relay and switching circuits”. *Electrical Engineering* 57.12 (1938), pp. 713–723 (cit. on p. 3).
- [SIMON 1988] Imre SIMON. “Recognizable sets with multiplicities in the tropical semiring”. In: *International Symposium on Mathematical Foundations of Computer Science*. Springer. 1988, pp. 107–120 (cit. on p. 8).
- [STAB and GUREVYCH 2017] Christian STAB and Iryna GUREVYCH. “Parsing argumentation structures in persuasive essays”. *Computational Linguistics* 43.3 (2017), pp. 619–659 (cit. on pp. 1, 27).
- [SYRJÄNEN and NIEMELÄ 2001] Tommi SYRJÄNEN and Ilkka NIEMELÄ. “The smodels system”. In: *International Conference on Logic Programming and NonMonotonic Reasoning*. Springer. 2001, pp. 434–438 (cit. on p. 21).
- [TONI, POTYKA, et al. 2023] Francesca TONI, Nico POTYKA, Markus ULRICHT, and Pietro TOTIS. “Understanding problog as probabilistic argumentation”. *arXiv preprint arXiv:2308.15891* (2023) (cit. on p. 28).
- [TONI and SERGOT 2011] Francesca TONI and Marek SERGOT. “Argumentation and answer set programming”. *Logic Programming, Knowledge Representation, and Non-monotonic Reasoning: Essays Dedicated to Michael Gelfond on the Occasion of His 65th Birthday* (2011), pp. 164–180 (cit. on pp. 13, 28).
- [TOTIS et al. 2023] Pietro TOTIS, Luc DE RAEDT, and Angelika KIMMIG. “Smproblog: stable model semantics in problog for probabilistic argumentation”. *Theory and Practice of Logic Programming* 23.6 (2023), pp. 1198–1247 (cit. on pp. 1, 28).
- [VALIANT 1979] I.g. VALIANT. “The complexity of computing the permanent”. *theoretical computer science* 8.2 (1979), pp. 189–201. ISSN: 0304-3975. DOI: [https://doi.org/10.1016/0304-3975\(79\)90044-6](https://doi.org/10.1016/0304-3975(79)90044-6). URL: <https://www.sciencedirect.com/science/Article/pii/0304397579900446> (cit. on p. 6).
- [VAN GELDER et al. 1991] Allen VAN GELDER, Kenneth A Ross, and John S SCHLIPF. “The well-founded semantics for general logic programs”. *Journal of the ACM (JACM)* 38.3 (1991), pp. 619–649 (cit. on p. 17).
- [WANG et al. 2025] Benjie WANG, Denis Deratani MAUÁ, Guy Van den BROECK, and YooJung CHOI. *A Compositional Atlas for Algebraic Circuits*. 2025. arXiv: 2412.05481 [cs.AI]. URL: <https://arxiv.org/abs/2412.05481> (cit. on pp. 30, 31).